LLMs Explained

FRONTMATTER

```
title: LLMs Explained
video no: 2 of 30
series: Building an AI Business From Scratch
voice_reference: ../../Input/4. Resources/2. scripting_resources/voice_reference/nate_herk_voice_reference.md
video_type: concept
hook_type: Curiosity-Gap
script_versions: A + B
status: DRAFT
```

REFERENCE VIDEOS

- Intro to Large Language Models (https://www.youtube.com/watch?v=zjkBMFhNj_g) -- Andrej Karpathy -- Watch for the "next word prediction" mental model and why emergent behaviour follows from it
- How Large Language Models Actually Work (Explained Simply) (https://www.youtube.com/watch?v=ceitP_wm79A) -- Independent creator -- Watch for the plain-English analogies used for non-technical audiences
- Claude Code Clearly Explained (https://www.youtube.com/watch?v=zxMjOqM7DFs) -- Independent creator -- Watch for how even a Claude-specific video has to explain LLMs first to frame everything else

THE ONE THING

Every AI tool you use -- ChatGPT, Claude, Gemini, Grok, DeepSeek -- runs on the same basic engine: a system trained to predict what word comes next, at a scale most people cannot imagine.

HOOK

You use AI every day -- but you have no idea what is actually happening when you type a message. Here is what ChatGPT, Claude, and Gemini actually are.

TARGET VIEWER

Someone who uses Claude or ChatGPT daily for work but could not explain what an LLM is if their boss asked.

PAYOFF

After watching, the viewer can explain in one sentence what an LLM is, why different models behave differently, and why this understanding makes them a better AI user -- starting today.

BEAT STRUCTURE

Beat -- The Black Box Problem

Point: Most AI users have no mental model of what is happening -- they see magic, and that makes them passive consumers instead of active operators.
Visual: Screen recording of typing into ChatGPT, response appears -- then cursor blinks on the question "what IS that?"
Target words: ~120w

Beat -- Prediction Machines

Point: LLMs are trained to predict the most likely next word, at massive scale, across billions of examples of human writing.
Visual: Simple text animation showing word-by-word prediction: "The sky is... [blue]" -- gradually building to paragraph-level output
Target words: ~200w

Beat -- Why They Are All Different

Point: Same basic architecture, different training data and tuning -- that is why Claude feels different from ChatGPT, and why Grok feels different from both.
Visual: Side-by-side: Claude, ChatGPT, Gemini given the same prompt -- highlight the tone and style differences
Target words: ~180w

Beat -- What This Means For You

Point: Understanding LLMs as prediction engines -- not search engines, not fact databases -- changes how you use them and why your prompts either work or do not.
Visual: Split screen: vague prompt vs structured prompt -- showing how framing the starting condition changes the prediction
Target words: ~180w

Beat -- CTA Bridge

Point: Now that you know what LLMs are, the next question is how to talk to them -- because most people are doing it completely wrong.
Visual: Preview clip from Prompts Explained video
Target words: ~60w

ANALOGY BANK

- What LLMs are: Autocomplete at civilizational scale -- like your phone autocomplete but trained on every book, article, and conversation ever written, predicting not just one word but entire coherent paragraphs
- Why different models differ: Same recipe, different kitchen -- every chef uses the same ingredients but results taste completely different based on technique, sourcing, and how long it was refined

THESIS LINE LOCATION

Beat 4 -- The insight is: LLMs are not searching for truth, they are predicting language -- which means your job is to give them the right starting conditions, not ask the right question.

FRAMEWORK PHRASE

Predict, not retrieve

SPECIFICITY BANK

- The five main LLMs: ChatGPT (OpenAI), Claude (Anthropic), Gemini (Google), Grok (xAI), DeepSeek (DeepSeek AI)
- Technical term for what they do: next token prediction
- Trained on trillions of tokens -- publicly acknowledged by all major labs
- Example models: claude-sonnet-4-6, GPT-4o, Gemini 2.5 Pro

CONSTRAINTS

- Do not explain transformer architecture or attention mechanisms
- Do not go into training process details (backpropagation, loss functions)
- Do not compare performance benchmarks -- that is for a different video
- Keep it conceptual -- this video is about understanding, not doing

CTA DIRECTION

Now that you know what LLMs are, the next video covers the only skill that immediately changes your results: prompting -- and most people are doing it completely wrong.

PERSONAL ANGLE [DANIEL TO FILL]

Suggested: The moment you realised you had been using AI like a search engine and started getting dramatically better results once you understood it as a prediction system -- when did that shift happen for you?

VULNERABILITY MOMENT [DANIEL TO FILL]

Suggested: You trusted a Claude answer completely early on -- it was confident, well-written, and wrong. The embarrassing moment where you learnt these are not fact databases.
Note: belongs in Beat 4