

Context Windows Explained

FRONTMATTER

title: Context Windows Explained
video no: 13 of 30
series: Building an AI Business From Scratch
voice_reference: ../../Input/4. Resources/2.
scripting_resources/voice_reference/nate_herk_voice_reference.md
video_type: concept
hook_type: Curiosity-Gap
script_versions: A + B
status: DRAFT

REFERENCE VIDEOS

- CLAUDE CODE FULL COURSE 4 HOURS: Build & Sell (2026) (<https://www.youtube.com/watch?v=QoQBzR1NIql>) -- Nick Saraev -- Watch the context management section: how Nick keeps Claude on track in long sessions and large projects
- Stop Vibe Coding: Context Engineering & RAG for AI Agents (<https://www.youtube.com/watch?v=7zZtxUW6dhc>) -- Cole Medin -- Watch for the distinction between context window limits and practical degradation -- performance drops before the limit hits
- What are LLM Tokens and API Prices? (Beginner Friendly) (<https://www.youtube.com/watch?v=7zbuEULTHIs>) -- Independent creator -- Watch as setup for the next video: tokens and context windows are related and this video bridges them

THE ONE THING

Claude does not forget -- it simply cannot see past a fixed boundary of text, and the practical solution is designing your work to stay inside that boundary.

HOOK

You are mid-conversation with Claude and it suddenly forgets what you said at the start. This is not a bug -- it is a hard physical limit. Here is how it works and what to do about it.

TARGET VIEWER

Someone who has experienced Claude "forgetting" earlier parts of a conversation or losing track of context in a long session and wants to understand why and how to prevent it.

PAYOFF

After watching, the viewer understands what a context window is, why it causes apparent "forgetting," and has three practical techniques to keep Claude accurate in long sessions and large projects.

BEAT STRUCTURE

Beat -- The Forgetting Moment

Point: You are in a long Claude session. It gives an answer that contradicts something it said an hour ago. This feels like Claude forgot -- but that is not quite what happened.

Visual: Screen recording of a long chat where Claude references something from early on correctly, then later contradicts it -- the exact moment the context window clipped it out

Target words: ~110w

Beat -- What a Context Window Actually Is

Point: A context window is the maximum amount of text Claude can see at once -- everything outside that window literally does not exist for Claude. It is not forgetting; it is a hard boundary.

Visual: Diagram: a sliding window moving along a long conversation -- text inside the window is lit, text outside is dark. As the conversation grows, early messages fall out.

Target words: ~180w

Beat -- Why Performance Degrades Before the Limit

Point: Claude does not suddenly fail when it hits the limit -- it gradually degrades. Information in the middle of long contexts gets less attention than information at the start or end.

Visual: Graph showing attention/performance on the Y axis vs position in context on the X axis -- showing the U-shaped curve: strong at start, weak in middle, stronger at end

Target words: ~160w

Beat -- Three Practical Workarounds

Point: Fresh session with a summary: start a new chat with a concise summary of what matters. CLAUDE.md for persistent context: put project facts there. Break tasks into smaller steps that fit cleanly.

Visual: Three mini-demos: (1) pasting a summary into a new session, (2) showing CLAUDE.md serving as persistent memory, (3) breaking a long task into two shorter ones

Target words: ~200w

Beat -- CTA Bridge

Point: Context windows are measured in tokens -- the next video explains what tokens are and how they connect to your subscription cost.

Visual: Preview clip from Tokens Explained

Target words: ~60w

ANALOGY BANK

- Context window: Like working memory for a human -- you can hold 7 plus or minus 2 things in mind at once. Ask someone to remember a list of 30 items and they will remember the first few, the last few, and forget the middle. Claude has the same constraint, just larger.
- Context degradation before limit: Like reading a very long document -- you remember the introduction and the conclusion, but the details in the middle get fuzzy. The context window has the same bias.

THESIS LINE LOCATION

Beat 2 -- The insight is: Claude is not forgetting -- it literally cannot see past a fixed boundary of text. Once you understand it as a viewing window rather than a memory system, the workarounds become obvious.

FRAMEWORK PHRASE

Window, not memory

SPECIFICITY BANK

- Claude claude-sonnet-4-6 context window: 200,000 tokens (approximately 150,000 words)
- Claude Opus 4.7: 200,000 token context window
- Approximately 750 words per 1,000 tokens (rough rule of thumb)
- Context compaction: Claude Code automatically summarises older context to stay within limits
- CLAUDE.md content persists through context compaction
- Performance degrades noticeably before hitting the hard limit -- practical limit is often 60-70% of the technical maximum

CONSTRAINTS

- Do not explain tokens in depth -- that is the next video
- Do not go into RAG or vector databases as solutions -- keep workarounds simple and native
- Keep it practical -- focus on understanding and workarounds, not technical architecture
- Do not explain how transformers work -- this is about practical impact, not theory

CTA DIRECTION

Context windows are measured in tokens. The next video explains what tokens are, how they connect to your AI subscription cost, and how to get more value from the same plan.

PERSONAL ANGLE [DANIEL TO FILL]

Suggested: A specific session where Claude lost important context and gave a contradictory answer -- and what you changed in your workflow to prevent it.

VULNERABILITY MOMENT [DANIEL TO FILL]

Suggested: A long Claude Code session where context loss caused Claude to undo work it had done earlier -- and the cost of not managing context actively.

Note: belongs in Beat 1