

Tokens Explained

FRONTMATTER

title: Tokens Explained
video no: 14 of 30
series: Building an AI Business From Scratch
voice_reference: ../../Input/4. Resources/2.
scripting_resources/voice_reference/nate_herk_voice_reference.md
video_type: concept
hook_type: Bold-Claim
script_versions: A + B
status: DRAFT

REFERENCE VIDEOS

- What are LLM Tokens and API Prices? (Beginner Friendly) (<https://www.youtube.com/watch?v=7zbuEULTHIs>) -- Independent creator -- Watch for the plain-English definition and the cost calculation walkthrough: this is the reference for accuracy
- CLAUDE CODE FULL COURSE 4 HOURS: Build & Sell (2026) (<https://www.youtube.com/watch?v=QoQBzR1NIql>) -- Nick Saraev -- Watch the token management section: how Nick reduces token usage without reducing output quality in real projects
- Stop Vibe Coding: Context Engineering & RAG for AI Agents (<https://www.youtube.com/watch?v=7zZtxUW6dhc>) -- Cole Medin -- Watch for the connection between context efficiency and token cost -- fewer tokens in means lower cost per task

THE ONE THING

Tokens are the currency of AI -- every word you send and receive costs tokens, and understanding that changes how you write prompts, structure workflows, and pick your subscription.

HOOK

You are paying for AI by the token whether you know it or not -- and most people are burning 30 percent more tokens than they need to. Here is what tokens are and how to spend them smarter.

TARGET VIEWER

Someone on a Claude Pro or Max subscription who has hit usage limits, seen API costs grow faster than expected, or simply never understood what they are actually paying for.

PAYOFF

After watching, the viewer understands what a token is, how input and output tokens differ in cost, and can apply three immediate changes to get more value from the same AI spend.

BEAT STRUCTURE

Beat -- The Hidden Currency

Point: Every time you send a message to Claude and get a response, tokens are consumed. Most users have no idea how many -- or how to reduce it without getting less out.

Visual: Screen: Claude session with a token counter visible, typing a prompt and watching tokens tick up -- then the response adding more

Target words: ~110w

Beat -- What a Token Actually Is

Point: A token is a chunk of text -- roughly a word, though common words split differently than uncommon ones. Around 750 words = 1,000 tokens. One page of text is approximately 500 tokens.

Visual: Animation: the sentence "The quick brown fox" breaking into tokens [The] [quick] [brown] [fox] -- then an unusual word breaking into multiple tokens

Target words: ~160w

Beat -- Input vs Output -- The Cost Difference

Point: Input tokens (what you send) cost less than output tokens (what Claude writes back). Most AI pricing has output at 3-5x the input price -- long Claude responses cost more than long Claude prompts.

Visual: Pricing table: Claude Sonnet input vs output cost per million tokens -- show the multiplier clearly

Target words: ~160w

Beat -- Three Ways to Get More Per Token

Point: Be specific not verbose in prompts (fewer input tokens). Ask for shorter outputs when you do not need long ones. Use cheaper models for simple tasks. These three changes reduce cost 20-40 percent with no loss.

Visual: Three before/after prompt examples: bloated vs tight. Show token count for each. Then model selection -- Haiku vs Sonnet for a simple task.

Target words: ~210w

Beat -- Subscription vs API -- Which to Use

Point: Claude Pro and Max subscriptions give you a monthly allocation. The API charges per token, pay-as-you-go. For regular daily use, subscription is cheaper. For automation, API is usually required.

Visual: Simple cost comparison: subscription monthly cost vs API cost at equivalent usage -- show the crossover point

Target words: ~130w

Beat -- CTA Bridge

Point: You now understand tokens. The next video covers the three ways to access AI programmatically: API, CLI, and MCP -- and when to use each.

Visual: Preview clip from API CLI MCP

Target words: ~60w

ANALOGY BANK

- Tokens: Like taxi fare by the word -- every word you say to the driver and every word they say back is on the meter. The longer the conversation, the higher the bill. Writing tight and purposeful is not just clarity, it is economy.

- Input vs output cost difference: Like a phone plan where incoming calls are cheap but outgoing calls are expensive -- or the reverse. Knowing which direction costs more changes how you design the conversation.

THESIS LINE LOCATION

Beat 1 -- The insight is: tokens are the currency of AI -- once you understand that every word has a cost, the way you write prompts, structure outputs, and pick models changes permanently.

FRAMEWORK PHRASE

Every word costs

SPECIFICITY BANK

- Rule of thumb: ~750 words per 1,000 tokens; ~500 tokens per page of text
- Claude claude-sonnet-4-6 pricing: input and output priced separately (check api.anthropic.com/pricing for current rates)
- Output tokens typically cost 3-5x input tokens across major providers
- Claude Haiku: fastest and cheapest; Claude Opus: most capable and most expensive
- Prompt caching: repeat context (like CLAUDE.md) can be cached to reduce input token cost on repeated calls

CONSTRAINTS

- Do not quote specific current prices -- they change; direct viewers to check the pricing page
- Do not go into token saving techniques specific to Claude Code -- that is a later video
- Do not explain context windows again -- that was the previous video; reference it but do not repeat it
- Keep it grounded in daily use and subscription context, not enterprise API budgeting

CTA DIRECTION

Now you understand tokens. The next video covers the three ways to access AI programmatically -- API, CLI, and MCP -- and when to use each one.

PERSONAL ANGLE [DANIEL TO FILL]

Suggested: The first time you saw an API bill that was higher than expected -- what caused it and what you changed immediately.

VULNERABILITY MOMENT [DANIEL TO FILL]

Suggested: Running a workflow that looped and consumed thousands of tokens before you caught it -- and how you added a cost check to the process after.

Note: belongs in Beat 1