Ollama Setup

FRONTMATTER

```
title: Ollama Setup
video no: 20 of 30
series: Building an AI Business From Scratch
voice_reference: ../../Input/4. Resources/2.
scripting_resources/voice_reference/nate_herk_voice_reference.md
video_type: tutorial
hook_type: Demonstration
script_versions: A + B
status: DRAFT
```

REFERENCE VIDEOS

- host ALL your AI locally (https://www.youtube.com/watch?v=Wjrdr0NU4Sk) -- NetworkChuck -- Watch for the energy and pacing of a technical setup video for a non-technical audience; notice how he handles the hardware requirements conversation
- Run AI Models Locally: Ollama Tutorial (Step-by-Step Guide + WebUI) (https://www.youtube.com/watch?v=Lb5D892-2HY) -- Independent creator -- Watch for the model selection section: how to explain the size vs capability tradeoff to beginners
- I'm changing how I use AI (Open WebUI + LiteLLM) (https://www.youtube.com/watch?v=nQCOTzS5oU0) -- NetworkChuck -- Watch for how local models integrate into a broader AI workflow -- this frames the next-video context well

THE ONE THING

Ollama is the easiest way to run AI models locally -- download it, pull a model with one command, and you have private AI running on your machine in under 10 minutes.

HOOK

What if you could run AI completely offline, for free, with no API key and no data sent anywhere? Ollama makes that possible. Here is the setup, start to finish.

TARGET VIEWER

Someone who uses Claude or ChatGPT via cloud and wants to experiment with running models locally -- motivated by privacy, cost reduction, or curiosity about offline AI.

PAYOFF

After watching, the viewer has Ollama installed, has pulled their first model, and has run a real prompt against it locally -- all completed in under 15 minutes from zero.

BEAT STRUCTURE

Beat -- Why Run AI Locally

Point: Three reasons: privacy (data never leaves your machine), cost (no API fees after download), and offline access (works anywhere, no internet required).

Visual: Three icons on screen: lock (privacy), dollar sign (zero per-query cost), wifi-off (offline capable). Then quick demo of Ollama running with wifi turned off.

Target words: ~100w

Beat -- What Ollama Is

Point: Ollama handles downloading, running, and managing AI models locally. It abstracts all the technical complexity -- you say "run llama3.2" and it handles the rest.

Visual: Browser: ollama.com model library showing 100+ models available. Then terminal: one command to pull and run. No config files, no setup beyond the download.

Target words: ~130w

Beat -- Installation -- One Download

Point: Go to ollama.com, download the installer for Mac, Windows, or Linux, run it. Done in under 2 minutes. Ollama runs as a background service after that.

Visual: Screen recording: ollama.com download page, installer running, Ollama appearing in system tray. Verify with ollama --version in terminal.

Target words: ~110w

Beat -- Pull and Run Your First Model

Point: Two commands: "ollama pull llama3.2" downloads the model. "ollama run llama3.2" starts it. Type your first prompt. Response appears locally with no internet.

Visual: Full terminal demo: ollama pull llama3.2 (download progress), ollama run llama3.2 (prompt appears), type a real question, response appears locally

Target words: ~200w

Beat -- Choosing the Right Starting Model

Point: Bigger models are smarter but need more RAM. For most machines: llama3.2:3b for speed, llama3.2:8b for balance. Check your RAM before pulling a 70B model.

Visual: Table: three model sizes (3B, 8B, 70B) -- RAM requirement, speed, capability indicator. Show where a typical 16GB MacBook sits.

Target words: ~160w

Beat -- CTA Bridge

Point: Ollama is running. Next: Local AI Explained -- when to use local vs cloud AI, and what running models locally actually unlocks for your workflow.

Visual: Preview clip from Local AI Explained

Target words: ~60w

ANALOGY BANK

- Ollama: Like Docker, but for AI models -- it handles the packaging, versioning, and running so you just interact with the interface, not the underlying infrastructure. If you know Docker, you already understand the concept.

- Choosing model size: Like choosing a car engine -- bigger is more powerful but uses more fuel and is harder to park. Match the engine to the use case, not to ego.

THESIS LINE LOCATION

Beat 4 -- The insight is: the local AI barrier is no longer technical -- Ollama reduced it to two terminal commands. The only remaining question is which model to start with.

FRAMEWORK PHRASE

Pull, run, done

SPECIFICITY BANK

- Install: download from ollama.com (Mac, Windows, Linux installers available)
- Pull command: ollama pull llama3.2
- Run command: ollama run llama3.2
- Check installed models: ollama list
- RAM requirements: 3B models need ~3GB, 8B needs ~8GB, 70B needs ~32GB+
- Available models: llama3.2, mistral, qwen2.5, deepseek-r1, phi4, gemma3, and 100+ at ollama.com/library
- Ollama exposes an OpenAI-compatible API at localhost:11434

CONSTRAINTS

- Do not go into Open WebUI, LiteLLM, or other Ollama frontends -- keep it terminal-only for simplicity
- Do not explain model quantisation -- keep it as a black box ("bigger needs more RAM")
- Do not do a detailed quality comparison with Claude or ChatGPT -- that is for later videos
- Primary target: Mac. Note Windows and Linux work the same way but do not spend time on them.

CTA DIRECTION

Ollama is running. The next video explains when to actually use local AI vs cloud AI -- and what running models locally unlocks that cloud AI structurally cannot give you.

PERSONAL ANGLE [DANIEL TO FILL]

Suggested: The first real task you used Ollama for instead of Claude -- what the task was and why local was the right choice for that specific use case.

VULNERABILITY MOMENT [DANIEL TO FILL]

Suggested: Expecting a local model to match Claude's quality and being caught off guard by the gap -- and how you recalibrated expectations to use local for the right tasks.
Note: belongs in Beat 5