

Local AI Explained

FRONTMATTER

title: Local AI Explained
video no: 21 of 30
series: Building an AI Business From Scratch
voice_reference: ../../Input/4. Resources/2.
scripting_resources/voice_reference/nate_herk_voice_reference.md
video_type: concept
hook_type: Curiosity-Gap
script_versions: A + B
status: DRAFT

REFERENCE VIDEOS

- host ALL your AI locally (<https://www.youtube.com/watch?v=Wjrdr0NU4Sk>) -- NetworkChuck -- Watch for how NetworkChuck frames the "why" of local AI for a non-technical audience: the emotional hook around privacy and ownership
- I'm changing how I use AI (Open WebUI + LiteLLM) (<https://www.youtube.com/watch?v=nQCOTzS5oU0>) -- NetworkChuck -- Watch for the hybrid workflow: local for some tasks, cloud for others -- this is the practical model this video should arrive at
- "I built 500+ AI agents, here's everything I know" (<https://www.youtube.com/watch?v=Svp7fbF0g2I>) -- Nate Herk -- Watch for how Nate thinks about model selection in production: matching model capability to task complexity

THE ONE THING

Local AI is not a budget version of cloud AI -- it is a different deployment model that solves problems cloud AI cannot, specifically around privacy, latency, and control.

HOOK

Everyone talks about local AI as a cheaper alternative to cloud AI. That is the wrong frame entirely. Here is what local AI actually gives you that cloud AI never will.

TARGET VIEWER

Someone who just set up Ollama and can run models locally but is not sure when to actually use them vs cloud AI -- they want a clear mental model for the decision.

PAYOFF

After watching, the viewer has a decision framework for when to use local vs cloud AI, understands the real tradeoffs on both sides, and can identify at least one current workflow where local is the better choice.

BEAT STRUCTURE

Beat -- The Wrong Frame

Point: Most people think of local AI as "cloud AI but cheaper." That reframe misses the point. Local AI solves different problems, not the same problem at lower cost.

Visual: Two columns on screen: "What people think local AI is" vs "What it actually is" -- each with 3 concrete bullet points. Pause on the contrast.

Target words: ~100w

Beat -- The Three Real Benefits

Point: Privacy (data never leaves your device), control (you choose the model, version, and configuration with no vendor dependency), and offline capability (works on a plane, in a restricted environment, without internet).

Visual: Three quick demos: 1) processing a sensitive PDF locally with no internet; 2) running a specific model version not available via API; 3) prompting a local model with wifi explicitly off

Target words: ~190w

Beat -- The Real Tradeoffs -- Be Honest

Point: Local models are smaller than frontier cloud models. Output quality on complex tasks is genuinely lower. Setup has a learning curve. Hardware limits what you can run. These are real constraints.

Visual: Honest side-by-side: llama3.2:8b vs claude-sonnet-4-6 on the same coding task -- show the output quality difference without softening it

Target words: ~170w

Beat -- The Decision Framework

Point: Use local for: sensitive data, offline environments, high-volume repetitive tasks where a lower quality threshold is acceptable, experimentation without API cost. Use cloud for: complex reasoning, creative work, anything where output quality is the primary constraint.

Visual: Decision flowchart: four yes/no questions pointing to local or cloud -- make it scannable and printable

Target words: ~190w

Beat -- Local Plus Cloud -- The Real Model

Point: For an AI business, local handles high-volume commodity tasks; cloud handles high-value complex work. Running both strategically is more efficient than committing to only one.

Visual: Workflow diagram: local model for first-pass processing and triage, cloud model for final quality work -- rough cost comparison per 1000 tasks

Target words: ~160w

Beat -- CTA Bridge

Point: You understand when to use local AI. Next: Codex Setup -- OpenAI's AI coding agent and how to add it to your toolkit alongside Claude Code.

Visual: Preview clip from Codex Setup

Target words: ~60w

ANALOGY BANK

- Local vs cloud AI: Like an in-house employee vs a remote specialist. The in-house person handles

sensitive documents you would never email to a contractor. The specialist is more capable but works through a third party. You need both for different jobs.

- Using local and cloud together: Like having a local pharmacist for everyday prescriptions and a specialist for complex diagnoses. Different capability levels for different task complexity levels -- the smart move is not picking one but knowing when to use each.

THESIS LINE LOCATION

Beat 4 -- The insight is: local AI is not a budget version of cloud AI -- it is the right tool for specific problems that cloud AI structurally cannot solve. The decision should be about fit, not cost.

FRAMEWORK PHRASE

Right model, right task

SPECIFICITY BANK

- Local models via Ollama: llama3.2, mistral, deepseek-r1, qwen2.5, phi4
- Claude claude-sonnet-4-6: cloud-only, 200K context, strongest reasoning available
- Quantised models: 70B+ models compressed to run on consumer hardware (typically 16-32GB RAM)
- Ollama OpenAI-compatible API at localhost:11434 -- drop-in for some integrations
- In 2026, quantised local models handle 95% of daily tasks acceptably (per local AI benchmarks)

CONSTRAINTS

- Do not go into technical integration with local APIs -- keep it conceptual
- Do not do a full model benchmark comparison -- that is for the Best LLMs video later in the series
- Be honest about quality tradeoffs -- do not position local as comparable to frontier cloud models on complex tasks
- Do not go into open-source model licences or legal considerations

CTA DIRECTION

You know when to use local AI. Next: Codex Setup -- OpenAI's AI coding tool and how to set it up alongside Claude Code in your workflow.

PERSONAL ANGLE [DANIEL TO FILL]

Suggested: A specific task you moved from cloud AI to local AI -- what was the trigger (privacy concern, cost, or offline need) and what you lost or gained.

VULNERABILITY MOMENT [DANIEL TO FILL]

Suggested: Using a local model for something it was not capable of and getting poor results that cost time to fix -- the specific lesson about matching model capability to task complexity.

Note: belongs in Beat 3