

25. Hermes Agents Explained

Frontmatter

video_number: 25
folder_title: Hermes_Agents_Explained
series: Building an AI Business From Scratch
video_type: Explainer / New Tool
hook_type: Contrast hook -- agents that forget vs agents that learn
framework_phrase: Learn, grow, run
script_versions: 1
status: ideation

Reference Videos

- The ONLY Hermes Agent Tutorial You Need 2026 --
<https://www.youtube.com/watch?v=r5FrGDTCSg>
- Full Hermes Agent Set-Up For Beginners in 2026! --
<https://www.youtube.com/watch?v=w4xOiuBQHKA>
- Hermes Agent Just Killed OpenClaw (Full Tutorial) -- https://www.youtube.com/watch?v=jmtpYUOr7_U

The One Thing

Hermes is the first agent that gets better at working with you over time -- not just better at tasks, but better at understanding how you work.

Hook

Every AI agent you've used forgets everything after the session ends. Hermes doesn't. It remembers, builds skills from experience, and runs while you sleep. Here's what that actually looks like.

Target Viewer

Builders who've maxed out what Claude Code can do in a session and want an agent that operates asynchronously, persistently, and without them being present.

Payoff

After watching, the viewer understands the three things that make Hermes different (persistent memory, self-improving skills, multi-platform access), knows what kind of tasks it's best suited for, and can describe it accurately to a client or teammate.

Beat Structure

Beat -- 1 -- The Memory Problem

Point: Every current AI agent resets at session end -- Claude Code, Codex, even custom GPTs. Hermes was built specifically to solve this: it creates and refines skills from every session.

Visual: Diagram: session ends, memory wipes vs Hermes session ends, skills updated

Target words: session memory, skill creation, persistent agent

Beat -- 2 -- What Hermes Actually Is

Point: An open-source agent by Nous Research that runs on your VPS, builds a skill library from experience, and learns your preferences across every interaction.

Visual: Hermes agent homepage screenshot, GitHub star count, Nous Research branding

Target words: Nous Research, open-source, VPS, self-hosted

Beat -- 3 -- The Skill Loop

Point: After each task, Hermes evaluates what it did, writes or updates a skill for that task type, and uses that skill next time -- getting faster and more accurate over repeated runs.

Visual: Diagram: task -> evaluation -> skill write -> skill applied next run

Target words: skill loop, self-improvement, task evaluation

Beat -- 4 -- Multi-Platform Access

Point: Hermes connects to Telegram, Discord, Slack, WhatsApp, Signal, and CLI through a single gateway -- you trigger tasks from wherever you are, it runs on your server.

Visual: Screen: one Hermes instance, multiple platform connections shown as icons

Target words: Telegram, Discord, gateway, multi-platform

Beat -- 5 -- First Setup and Use Case Fit

Point: Setup is a 45-minute VPS deploy. Best fit: async knowledge work, recurring research tasks, and workflows where the agent improving over time has compounding value.

Visual: Terminal: Hermes install commands, first task via Telegram

Target words: VPS deploy, async, recurring tasks, compounding value

Analogy Bank

- Hermes vs standard agents: like the difference between a contractor who starts fresh each project vs a long-term employee who knows your systems and gets faster every month. [concept: persistent memory value]
- The skill loop: like a chef who writes down what worked after each service and refines the recipe -- by month 3 the dish is better than anything you wrote in the original brief. [concept: self-improvement mechanism]

Thesis Line Location

Beat 3 -- the skill loop is the core insight: Hermes doesn't just remember, it actively improves the way it performs recurring tasks, which compounds in a way no session-based agent can match.

Framework Phrase

Learn, grow, run

Specificity Bank

- Hermes Agent -- built by Nous Research, open-source on GitHub

- 100k+ GitHub stars (as of May 2026)
- Runs on VPS -- setup time approximately 45 minutes
- Platforms: Telegram, Discord, Slack, WhatsApp, Signal, CLI
- Nate Herk described it as A-tier in his weekly stack (confirmed via X post)

Constraints

- Do not compare Hermes vs OpenClaw directly here -- that's Video 26's territory
- Do not oversell persistent memory -- note it's only as good as the skill quality it writes
- Keep setup steps high-level -- a full tutorial would be its own video

CTA Direction

Next video: Hermes isn't the only persistent agent in this space -- OpenClaw takes a different approach and the difference matters depending on what you're building.

Personal Angle

[DANIEL TO FILL] -- e.g. the first task you gave Hermes that you expected to explain twice and it handled correctly on the second run without any new instructions

Vulnerability Moment

[DANIEL TO FILL] -- e.g. the skill it wrote from a bad session that you didn't catch until it used it incorrectly three times in a row