29. Best LLMs

Frontmatter

video_number: 29
folder_title: Best_LLMs
series: Building an AI Business From Scratch
video_type: Decision Guide / Model Comparison
hook_type: Reframe hook -- wrong question, right question
framework_phrase: Task, cost, fit
script_versions: 1
status: ideation

Reference Videos

- AI Model Comparison 2026: ChatGPT vs Claude, Perplexity & More -- https://www.youtube.com/watch?v=FbBNLYw_dRE
- LLMs in 2026: Trends, Tools, & How to Choose the Right One -- https://www.youtube.com/watch?v=nPFT5iVFktg
- The Best LLM Is.... (A breakdown for every category) -- https://www.youtube.com/watch?v=0K66T6J1pVc

The One Thing

There's no best LLM -- there's the best LLM for each task type, and knowing the difference saves money and improves output at the same time.

Hook

Everyone asks which LLM is best. That's the wrong question. The right question is: best for what? Here's the framework I use -- and my current picks for every task type I actually work on.

Target Viewer

Builders who need to make real model decisions for client work or internal systems -- someone who's been using Claude Code and now wants to know when to reach for a different model.

Payoff

After watching, the viewer knows the four task categories and which model wins each one in 2026, can make a model choice in under 30 seconds using the framework, and understands the cost-vs-capability tradeoff for each tier.

Beat Structure

Beat -- 1 -- Why Benchmarks Mislead

Point: Benchmark scores measure performance on controlled tasks -- most real work is messier, more

context-dependent, and less about raw capability than about fit with your workflow and prompt style.
Visual: Benchmark leaderboard screenshot vs 'tasks I actually do' list -- almost no overlap
Target words: benchmark, controlled task, real-world fit, misleading

Beat -- 2 -- The Four Task Categories

Point: Coding (needs instruction-following, file-level reasoning). Reasoning (needs chain-of-thought, planning). Writing (needs voice, tone control, creativity). Retrieval (needs context handling, accuracy over style).
Visual: Four-quadrant grid: task category + what matters most for each
Target words: coding, reasoning, writing, retrieval, task fit

Beat -- 3 -- Best Models by Category (May 2026)

Point: Coding: Claude Sonnet 4.5 (best instruction-following) or Codex (o3, best for isolated reasoning tasks). Reasoning: Claude Opus 4.7 or o3. Writing: Claude Sonnet 4.5. Retrieval/RAG: Gemini 2.5 Pro (long context strength).
Visual: Model recommendation card per category: model name, why, when to upgrade or downgrade
Target words: Claude Sonnet 4.5, Claude Opus 4.7, o3, Gemini 2.5 Pro, category fit

Beat -- 4 -- Cost Considerations

Point: Cost scales with context and complexity. Use Haiku for classification, routing, and simple extraction. Use Sonnet for most agentic and coding work. Reserve Opus for high-stakes reasoning tasks where quality difference justifies 5-10x price.
Visual: Cost-tier table: Haiku/Sonnet/Opus with price per million tokens and task fit
Target words: Haiku, Sonnet, Opus, cost tier, task routing

Beat -- 5 -- How Model Choice Changes at Scale

Point: When a workflow runs 1,000 times a day, a 10x cost difference between models becomes the biggest line item. Model selection is infrastructure at scale -- optimize late, not early.
Visual: Cost projection chart: 100 runs/day vs 1,000 runs/day for Sonnet vs Haiku on same task
Target words: scale, cost projection, optimization timing, infrastructure

Analogy Bank

- Benchmarks vs task fit: like hiring a chef based on their knife skills test score -- it measures something real, but it doesn't tell you if they can run your kitchen. [concept: why benchmarks mislead]
- Model tiers: like car engine sizes -- a V8 is not always better than a V4, it depends on whether you're towing or commuting. Bigger model, bigger cost, not always better output. [concept: matching model to task]

Thesis Line Location

Beat 2 -- the four task categories are the framework; everything else (model picks, cost tiers, scale considerations) is application of the framework.

Framework Phrase

Task, cost, fit

Specificity Bank

- Claude Sonnet 4.5 -- best for coding and writing tasks (May 2026)
- Claude Opus 4.7 -- best for high-stakes reasoning; ~5-10x more expensive than Sonnet
- Claude Haiku 4.5 -- classification, routing, extraction; lowest cost tier
- o3 (Codex CLI default) -- best for isolated reasoning tasks in coding
- Gemini 2.5 Pro -- strongest long-context retrieval, 1M token window

Constraints

- This video is a snapshot in time -- acknowledge model recommendations change; the framework is permanent, the picks aren't
- Do not rate models on benchmarks alone -- always tie to a real task type
- Keep to models Daniel has personally used for work, not everything on the leaderboard

CTA Direction

Next video: now you know which model to use -- here's how the best builders are wiring these models together into agent frameworks, and which framework actually fits which type of work.

Personal Angle

[DANIEL TO FILL] -- e.g. the task where you switched from Sonnet to Opus and the quality difference justified the cost, and the task where you moved from Opus to Haiku and nobody noticed

Vulnerability Moment

[DANIEL TO FILL] -- e.g. the month where your Opus usage was 80% of your AI bill for tasks that Sonnet would have handled fine -- what that cost you and when you caught it