

31. AI Coding Setup

Frontmatter

video_number: 31
folder_title: AI_Coding_Setup
series: Building an AI Business From Scratch
video_type: Series Finale / Full Stack Reveal
hook_type: Authority hook -- here's what I actually use
framework_phrase: Model, environment, workflow
script_versions: 1
status: ideation

Reference Videos

- I Tried Every AI Coding Agent... Here's My 2026 Setup --
<https://www.youtube.com/watch?v=zgxorh9LhiE>
- The Best AI Coding Tools in 2026 (That I Actually Use) --
<https://www.youtube.com/watch?v=xpFnbDPtTNI>
- Best AI Coding Tools for Developers in 2026 (Don't Choose Wrong) --
<https://www.youtube.com/watch?v=RKbmqSRc0z0>

The One Thing

Your AI coding setup isn't one tool -- it's a stack of compounding decisions, and most people optimize the wrong layer first.

Hook

I've tested every AI coding agent. Here's the setup I actually use every day, what I dropped, the one thing most setups get wrong, and how to build yours from the tools in this series.

Target Viewer

Someone who's been through the full series and wants to see how everything fits together into a real daily workflow -- or a new viewer who wants a practitioner's honest stack reveal.

Payoff

After watching, the viewer has a clear picture of a real, proven AI coding stack, knows the one configuration decision that has the biggest ROI, and can assemble their own version of the stack in one afternoon.

Beat Structure

Beat -- 1 -- What a Full Stack Looks Like

Point: The stack has four layers: model (which AI), environment (IDE + terminal), context (CLAUDE.md + project structure), and workflow (how tasks flow through the system). Most people only think about

layer one.

Visual: Four-layer stack diagram: model -> environment -> context -> workflow

Target words: four layers, model, environment, context, workflow

Beat -- 2 -- The Model Layer

Point: Claude Sonnet 4.5 as daily driver. Claude Opus 4.7 for complex architectural decisions. Haiku for fast classification tasks in automated pipelines. Codex CLI for isolated reasoning tasks where o3's depth matters.

Visual: Model selection card: each model, when it appears in the workflow, cost tier

Target words: Sonnet 4.5, Opus 4.7, Haiku, Codex, model routing

Beat -- 3 -- The Environment Layer

Point: VS Code with the Claude Code extension for in-editor context. Claude Code terminal as the primary agent interface for complex, multi-step tasks. The terminal is the power layer -- IDE plugin is the convenience layer.

Visual: Screen: VS Code with Claude extension + separate terminal with Claude Code running in parallel

Target words: VS Code, Claude extension, terminal agent, primary interface

Beat -- 4 -- The Context Layer (Most Underrated)

Point: CLAUDE.md is the highest-ROI configuration in the whole stack. A well-designed CLAUDE.md reduces back-and-forth by 60%, cuts token usage, and makes Claude work the way you work -- not the way it defaults.

Visual: Before/after session: same task with bare CLAUDE.md vs optimized CLAUDE.md, turn count and token comparison

Target words: CLAUDE.md, context layer, highest ROI, session efficiency

Beat -- 5 -- The Workflow Layer

Point: Tasks route by type: Claude Code for in-session coding. Hermes for async recurring tasks. Codex for well-scoped isolated jobs. n8n for stable webhook integrations. The workflow layer decides which tool handles which task.

Visual: Workflow routing diagram: task type -> tool assignment -> output destination

Target words: workflow routing, task type, tool assignment, full stack integration

Beat -- 6 -- Building Your Version

Point: Start with just Claude Code + a CLAUDE.md. Add Hermes when you have 3+ recurring tasks that take more than an hour each. Add Codex when you have isolated tasks where o3's reasoning is measurably better. Add n8n when you have stable, webhook-driven integrations that won't change for months.

Visual: Timeline: week 1 setup -> month 1 additions -> month 3 full stack, each with clear trigger condition

Target words: build in stages, trigger condition, week 1, month 1, month 3

Analogy Bank

- The four-layer stack: like a kitchen setup -- you need the ingredients (model), the equipment

(environment), the recipe (context/CLAUDE.md), and the workflow (how dishes move from prep to plate). Most people just buy a new knife. [concept: what a full stack means]

- CLAUDE.md as the highest ROI layer: like a surgeon's preference card -- 2 pages that mean every instrument, every preference, every constraint is ready before the operation starts. The card costs nothing; wrong setup costs hours. [concept: context layer value]

Thesis Line Location

Beat 4 -- CLAUDE.md as the highest-ROI layer is the core insight: most people optimize the model when they should optimize the context, and the gap between a bare setup and an optimized one is the biggest performance difference in the whole stack.

Framework Phrase

Model, environment, workflow

Specificity Bank

- Claude Sonnet 4.5 -- daily driver, best cost-to-performance for agentic coding (May 2026)
- Claude Opus 4.7 -- architectural decisions, complex reasoning; ~5-10x Sonnet cost
- Haiku 4.5 -- classification, routing; lowest cost tier
- CLAUDE.md optimization: 60% reduction in back-and-forth turns (observed across builders)
- Setup sequence: Claude Code + CLAUDE.md first -> Hermes at 3+ recurring tasks -> Codex for isolated reasoning -> n8n for stable webhooks

Constraints

- Do not make this a tutorial for any single tool -- this is a synthesis video, each tool has its own video in the series
- Do not recommend building the full stack on day one -- the staged approach is the point
- This is the series finale -- acknowledge the journey, but keep it practical, not sentimental

CTA Direction

Series complete. The next step is the Skool community -- where we build these setups together, share what's working, and go deeper on implementation than any YouTube video can.

Personal Angle

[DANIEL TO FILL] -- e.g. what your setup looked like on day 1 of the series vs what it looks like now, and the single change that had the most impact

Vulnerability Moment

[DANIEL TO FILL] -- e.g. the month where you had all the tools but no coherent workflow and felt like the stack was fighting you instead of helping you -- and what clicked to fix it